

Flow-based Network Intrusion Detection Based on BERT Masked Language Model

Loc Gia Nguyen
Nagaoka University of Technology
Nagaoka, Niigata, Japan
s203145@stn.nagaokaut.ac.jp

Kohei Watabe
Nagaoka University of Technology
Nagaoka, Niigata, Japan
k_watabe@vos.nagaokaut.ac.jp

ABSTRACT

A Network Intrusion Detection System (NIDS) is an important tool that identifies potential threats to a network. Recently, different flow-based NIDS designs utilizing Machine Learning (ML) algorithms have been proposed as potential solutions to detect intrusions efficiently. However, conventional ML-based classifiers have not seen widespread adoption in the real-world due to their poor domain adaptation capability. In this research, our goal is to explore the possibility of improve the domain adaptation capability of NIDS. Our proposal employs Natural Language Processing (NLP) techniques and Bidirectional Encoder Representations from Transformers (BERT) framework. The proposed method achieved positive results when tested on data from different domains.

CCS CONCEPTS

• Security and privacy → Intrusion detection systems.

KEYWORDS

Intrusion Detection System, Domain adaptation, BERT

ACM Reference Format:

Loc Gia Nguyen and Kohei Watabe. 2022. Flow-based Network Intrusion Detection Based on BERT Masked Language Model. In *CoNEXT Student Workshop 2022 (CoNEXT-SW '22)*, December 9, 2022, Roma, Italy. ACM, New York, NY, USA, 2 pages. <https://doi.org/10.1145/3565477.3569152>

1 INTRODUCTION

It is common in practical application of NIDS for there to be a change in the data distribution between its training data and the data it encounters when deployed. Conventional ML algorithms often adapt poorly to such change, which limit their usefulness in real-world scenarios [1]. To address this, Energy-based Flow Classifier (EFC) [3] was proposed as a solution. Despite having good adaptability, EFC produces high false positives rate for domains where the distribution of features of malicious flows overlap with that of benign flows. We theorize that the reason for the limitations of conventional ML algorithms and EFC is the use of singular flows as input data, as the classifier can only model the distribution of features within a flow. This limitation can be overcome with the use of sequences of flows, allowing the classifier to further models




Figure 1: Proposed system architecture

the distribution of a flow in relation to other flows. To utilize the context information from a sequence of flow, we use the BERT framework, which is able to process inputs in relation to all the other inputs in a sequence.

BERT[2] is a transformer-based machine learning technique for NLP developed by Google. The BERT framework is comprised of two steps: pre-training and fine-tuning. In pre-training, the BERT model is trained on unlabeled data. For fine-tuning, the model is first initialized using the pre-trained parameters, and then trained using labeled data from the downstream tasks. BERT is pre-trained with two unsupervised tasks, which are Masked Language Modeling (MLM) and Next Sentence Prediction (NSP). In MLM, some of the words in a sentence are replaced with a different token. The objective is to predict the original value of the masked words based on other unmasked words in the sentence. In NSP, BERT takes sentence pairs as input. The objective is to predict whether the second sentence in the pair is the next sentence in the document. For fine-tuning, task-specific inputs and outputs are added to a pre-trained BERT model.

Our research employs CIDDs-001[4] and CIDDs-002[5] data sets that contains flow samples from a small business environment emulated using OpenStack. CIDDs-001 also contains real traffic flow samples captured from an external server directly deployed on the internet.

2 PROPOSAL

We first organize network traffic flows into structures similar to a language, treating a flow as a word and a sequence of flows as a sentence. In this study, the BERT model is pre-trained with only the MLM task. For fine-tuning, a linear layer with softmax output is used. It is important to preserve the distribution of flows within a sequence; therefore, during training, the data set is not shuffled. A training sample is generated by selecting a segment of flows from the data set at random.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).
CoNEXT-SW '22, December 9, 2022, Roma, Italy
© 2022 Copyright held by the owner/author(s).
ACM ISBN 978-1-4503-9937-1/22/12.
<https://doi.org/10.1145/3565477.3569152>

Table 1: Average composition of data sets

CIDS-001 OpenStack		CIDS-001 External Server		CIDS-002	
1 set		10 sets		10 sets	
CIDS-001 large		CIDS-001 internal		CIDS-001 external	
label	#	label	#	label	#
<i>normal</i>	25178585	<i>normal</i>	10000	<i>unknown</i>	10000
<i>dos</i>	2775655	<i>dos</i>	9000	<i>suspicious</i>	10000
<i>portScan</i>	194642	<i>portScan</i>	935		
<i>pingScan</i>	5266	<i>pingScan</i>	45		
<i>bruteForce</i>	4992	<i>bruteForce</i>	20		
Total	28159140	Total	20000	Total	20000

The overall architecture of the system is illustrated in Figure 1. Six features from each flow (Src Pt, Dst Pt, Proto+Flags, Packets, Bytes, Duration) are used as input data, these features are discretized as described in EFC’s original paper [3]. The discrete value of each features are encoded as numbers ($flow_i$). BERT decodes each number into a 128-dimension vector, concatenates them to form a 768-dimension vector (e_{flow_i}), and processes it to produce a different 768-dimension vector (h_{flow_i}). The output of BERT is then passed through a Multilayer Perceptron classifier (a linear layer with softmax output), which reduces the dimension from 768 to 2. This 2-dimension vector represents the predicted probability of the flow being benign and malicious (e.g. benign 0.99, malicious 0.01). The predicted class of the flow is the class with the higher probability (e.g. benign).

We used data from three different domains: CIDS-001 OpenStack, CIDS-001 External Server, and CIDS-002 to evaluate the domain adaptation capability of the proposed method. Average composition of the data sets used in the experiment are shown in Table 1. Training was performed on one set of *CIDS-001 large*. While testing was performed on *CIDS-001 internal*, *CIDS-001 external*, and *CIDS-002*, each containing ten sets randomly selected from the full data sets. This testing scheme mimics the one used by Camila *et al.* [3] to make the results of the proposed method more comparable to those of EFC. Flows labeled *normal* are considered benign, while those labeled otherwise are considered malicious. For *CIDS-001 external*, flows labeled *unknown* and *suspicious* are considered benign and malicious respectively.

We assess the performance of our method in comparison to EFC and ML classifiers including Decision Tree (DT), K-Nearest Neighbors (KNN), Multilayer Perceptron (MLP), Naive Bayes (NB), and Support Vector Machine (SVM). Each classifier’s performance is measured using Accuracy and F1 score [1].

3 RESULTS AND DISCUSSION

Table 2 shows the average performance and standard error for each classifier. All classifiers achieved higher Accuracy and F1 Score on *CIDS-001 internal* test sets (same domain as training data) compared to the other test sets (different domains from training data). Both the proposed method and EFC maintained performance across the two different domains, with the proposed method outperforming EFC.

We also experimented with training the classifiers on smaller but balanced data sets (containing 80000 flows with the same proportion of labels as in *CIDS-001 internal*). However, performance was

Table 2: Average classification performance and standard error

Classifier	Train CIDS-001 large					
	Test CIDS-001 internal		Test CIDS-001 external		Test CIDS-002	
	Accuracy	F1 score	Accuracy	F1 score	Accuracy	F1 score
Proposal	0.994(0.002)	0.994(0.002)	0.895(0.027)	0.904(0.022)	0.916(0.045)	0.877(0.072)
EFC	0.941(0.005)	0.941(0.004)	0.747(0.039)	0.796(0.025)	0.846(0.046)	0.800(0.070)
DT	0.996(0.001)	0.996(0.001)	0.870(0.028)	0.874(0.024)	0.818(0.061)	0.707(0.106)
KNN	0.989(0.002)	0.989(0.003)	0.839(0.008)	0.811(0.011)	0.818(0.061)	0.707(0.106)
MLP	0.992(0.001)	0.992(0.001)	0.573(0.009)	0.285(0.023)	0.832(0.059)	0.729(0.108)
NB	0.903(0.002)	0.892(0.002)	0.500(0.000)	0.000(0.000)	0.499(0.000)	0.001(0.000)
SVM	0.570(0.015)	0.245(0.047)	0.738(0.016)	0.718(0.021)	0.513(0.013)	0.090(0.037)

worse for all classifiers when compared to those trained on *CIDS-001 large*. Notably, the performance of the proposed method was significantly affected for all domains. By creating balanced data sets, the distribution of flows within a sequence was also altered. This suggests that the distribution of flows within a sequence is learned by the model of the proposed method.

4 CONCLUSION AND FUTURE WORK

In this study, we suggest the use of singular flows input to be a possible explanation for the poor domain adaptation capability of conventional ML-based classifiers. Then we proposed the used of sequences of flows to address this limitation. We utilized BERT model for the representation of flow sequences and an MLP classifier to discriminate between benign and malicious flows. Early experimental results showed that the proposed method is capable of achieving good and consistent results across different domains. However, more extensive testing on recent data sets is needed to further evaluate its domain adaptation capability.

In future work, we plan to investigate the impact of flow sequence sampling method on result, such as using only benign flows or grouping flows into sequences that originate from the same hosts. Making the system less reliant on labeled data is another research goals of ours. We aim to achieve this by modeling sequences of benign flows then look for anomalous representations produce by BERT, indicating malicious flows.

ACKNOWLEDGMENTS

This work was partly supported by JSPS KAKENHI Grant Number JP20H04172.

REFERENCES

- [1] Zeeshan Ahmad, Adnan Shahid Khan, Cheah Wai Shiang, Johari Abdullah, and Farhan Ahmad. 2021. Network Intrusion Detection System: A Systematic Study of Machine Learning and Deep Learning Approaches. *Trans. Emerg. Telecommun. Technol.* 32, 1 (jan 2021), 30 pages. <https://doi.org/10.1002/ett.4150>
- [2] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Association for Computational Linguistics, Minneapolis, Minnesota, 4171–4186. <https://doi.org/10.18653/v1/N19-1423>
- [3] Camila F. T. Pontes, Manuela M. C. de Souza, João J. C. Gondim, Matt Bishop, and Marcelo Antonio Marotta. 2021. A New Method for Flow-Based Network Intrusion Detection Using the Inverse Potts Model. *IEEE Transactions on Network and Service Management* 18, 2 (2021), 1125–1136. <https://doi.org/10.1109/TNSM.2021.3075503>
- [4] Markus Ring, Sarah Wunderlich, Dominik Grödl, Dieter Landes, and Andreas Hotho. 2017. Flow-based benchmark data sets for intrusion detection. In *Proceedings of the 16th European Conference on Cyber Warfare and Security*. 361–369.
- [5] Markus Ring, Sarah Wunderlich, Dominik Grödl, Dieter Landes, and Andreas Hotho. 2017. Creation of Flow-Based Data Sets for Intrusion Detection. *Journal of Information Warfare* 16 (12 2017), 41–54.